

CS 344 : Operating Systems Laboratory

Lab 4

Members~

<i>Harsh</i>	<i>200123022</i>
<i>Pragyan Banerjee</i>	<i>200123080</i>
<i>Sahil Kumar Gupta</i>	<i>200123081</i>

PART A

What is lazy memory allocation?

Lazy allocation of memory means not allocating memory to a process until it is actually needed. By delaying allocation of memory until you actually need it, we can decrease startup time, and even eliminate the allocation entirely if we never actually use the process. Most modern operating systems perform lazy allocation of heap memory, though **xv6 does not**. Xv6 applications ask the kernel for heap memory using the `sbrk()` system call. In the original xv6 kernel, **`sbrk()` allocates physical memory and maps it into the process's virtual address space.**

Task: We have to add support for this lazy allocation feature in xv6 by delaying the memory requested by `sbrk()` until the process actually uses it.

Step 1: Eliminate allocation from `sbrk()`:

To delete the page allocation from the `sbrk(n)` system call implementation, which is in the function `sys sbrk()` in `sysproc.c`, we have used the given patch file in which the call to the function `growproc(n)` is commented out. The new `sbrk(n)` will just increment the process's size by `n` and return the old size. It does not allocate memory. However, it still increases the size of the process by `n` bytes to trick the process into believing that it has the memory requested.

Typing `echo hi` to the shell gives us the following output:

```
objdump -t kernel | sed '1,/SYMBOL TABLE/d; s/ .* / /; /^$/d' > kernel.sym
dd if=/dev/zero of=xv6.img count=10000
10000+0 records in
10000+0 records out
5120000 bytes (5.1 MB, 4.9 MiB) copied, 0.0201198 s, 254 MB/s
dd if=bootblock of=xv6.img conv=notrunc
1+0 records in
1+0 records out
512 bytes copied, 6.5864e-05 s, 7.8 MB/s
dd if=kernel of=xv6.img seek=1 conv=notrunc
349+1 records in
349+1 records out
178728 bytes (179 kB, 175 KiB) copied, 0.000593334 s, 301 MB/s
qemu-system-i386 -nographic -drive file=fs.img,index=1,media=disk,format=raw -drive file=xv6.img,index=0,media=disk,format=raw -smp 2 -m 512
xv6...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ echo hi
pid 3 sh: trap 14 err 6 on cpu 0 eip 0x112c addr 0x4004--kill proc
$
```

The “pid 3 sh: trap...” message is from the kernel trap handler in `trap.c`; it has caught a page fault (trap 14, or `T_PGFLT`), which the `xv6` kernel does not know how to handle. The “addr 0x4004” indicates that the virtual address that caused the page fault is 0x4004.

Step 2: Adding Lazy Allocation

We have modified the code of `trap()` function in `trap.c` to respond to a page fault from user space by mapping a newly-allocated page of physical memory at the faulting address, and then returning back to user space to let the process continue executing. We have allocated a new memory page, added suitable page table entries, and returned from the trap, so that the process can avoid the page fault the next time it runs.

Commented out portion of `sbrk` in `sysproc.c` provided in patch

```
44     int
45     sys_sbrk(void)
46     {
47         int addr;
48         int n;
49
50         if(argint(0, &n) < 0)
51             return -1;
52         addr = myproc()->sz;
53         myproc()->sz += n;
54
55         // if(growproc(n) < 0)
56         //     return -1;
57         return addr;
58     }
```

We have added the following piece of code in trap.c

Now we can begin to handle this page fault error starting from trap.c. In trap.c, there are various cases to handle trap errors using case-switch statements so we add one more for T_PGFLT since we have a page fault error and there is no pre-existing case for handling page-fault errors. To make mappages function available in trap.c we made a new header file named vm.h and included it in trap.c

vm.h

```
1 int mappages(pde_t *, void *, uint, uint, int);
```

Mappages function

```
int
mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm)
{
    char *a, *last;
    pte_t *pte;

    a = (char*)PGROUNDDOWN((uint)va);
    last = (char*)PGROUNDDOWN(((uint)va) + size - 1);
    for(;;){
        if((pte = walkpgdir(pgdir, a, 1)) == 0)
            return -1;
        if(*pte & PTE_P)
            panic("remap");
        *pte = pa | perm | PTE_P;
        if(a == last)
            break;
        a += PGSIZE;
        pa += PGSIZE;
    }
    return 0;
}
```

Changes in trap.c trap function

```
37 void
38 trap(struct trapframe *tf)
39 {
40     char *mem;
41     uint a;
```

```

83  case T_PGFLT:
84      mem = kalloc();
85      a = rcr2();
86      a = PGROUNDDOWN(a);
87      memset(mem, 0, PGSIZE);
88      mappages(myproc()->pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U);
89      break;

```

Important Points:-

1. `tf->trapno = T_PGFLT` check whether a fault is a page fault
2. `PGROUNDDOWN(rcr2())` is used to round the faulting virtual address down to start of page
3. We have returned in order to avoid the `cprintf()` and the `proc->killed = 1` statements
4. We have made a `vm.h` to access `mappages` in `trap.c`

Description of functions we changed and why we changed and what effect it had

PGROUNDUP(address) which rounds up the address to the starting of a page or we can say it gives us the page-aligned address since we can have errors at any address in-btw page.

kalloc() allocates one-page of 4096 bytes from physical memory & returns a pointer to that page.

memset() makes the whole page null.

mappages() maps the page to our process's page directory by converting the given virtual address (rounded) to physical address by the use of `V2P(mem)` which basically subtracts the Kernel base address from the virtual address. It creates a new (page table entry) PTE for that particular virtual address.

We use these functions in our lazy allocation implementation by giving the argument `myproc()` which contains the current process's directory and `rcr2()` which gives the address at which page fault occurs and the rest of the implementation is done in the `lazy_allocation` function. We need not use the old size and new size arguments since we only need to allocate a page so that our program can execute normally in userspace.

The page table flags used in `mappages()` function are `PTE_W` or `PTE_U` which declares that the page is user accessible and writeable as seen below.

Hence our new function takes the faulty address, rounds up and makes it page aligned, allocates a fresh page from the physical memory to the process using `mappages()` function and the process starts to execute normally if the allocation is done correctly.

After making the following changes, the lazy memory allocation is implemented correctly in xv6 and now we are not encountering any errors after typing `echo hi` OR `ls` to the shell as shown in the screenshot attached below:

```
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ echo hi
hi
```

PART B

Answers to Questions

1. How does the kernel know which physical pages are used and unused?

```
21 struct {
22     struct spinlock lock;
23     int use_lock;
24     struct run *freelist;
25 } kmem;
```

In `kalloc.c`, xv6 maintains a linked list of unoccupied pages under the name `kmem`. When the list is initially empty, xv6 runs `kinit1` through `main()`, adding 4MB of free pages to the list.

2. What data structures are used to answer this question?

`freelist`, a linked list as seen in the image above. The linked list's nodes are all instances of the `struct run`, which is defined in `kalloc.c` (pages are typecast to `(struct run *)` when inserted into the `freelist` in `kfree(char *v)`).

3. Where do these reside?

Within the structure `kmem` of the file `kalloc.c`, this linked list is defined. Every node is of the `struct run` type, which is specified in `kalloc.c` as well.

4. Does xv6 memory mechanism limit the number of user processes?

The number of user processes is constrained in xv6 due to a restriction on the size of `ptable` (a maximum of `NPROC` elements which is set to 6\$ by default). The `param.h` file defines `NPROC`.

5. If so, what is the lowest number of processes xv6 can 'have' at the same time (assuming the kernel requires no memory whatsoever)?

There is just one process running when the xv6 operating system boots up, and its name is `initproc` (this process forks the `sh` process which forks other user processes). Additionally, since a process can use up to 2GB (`KERNBASE`) of virtual address space and 240 MB (`PHYSTOP`) of maximum physical memory, one process can use up the whole physical memory (We added this since the question asks from a memory management perspective). So, the response is 1. Additionally, since all user interactions must be performed through user processes that are forked from `initproc/sh`, there cannot be zero processes after boot.

Task 1

In `proc.c`, the `create_kernel_process()` method was developed. The kernel process will continue to operate exclusively in kernel mode. Therefore, there is no need to initialise its `trapframe`, user space, or the user sector of its page table (`trapframes` hold userspace register values). The address of the next instruction is kept in the process' context's `eip` register. At the entrance point, we want the process to begin executing (which is a function pointer).

Consequently, we set the context's `eip` value to entry point (Since entry point is the address of a function). The process is given a place in the `ptable` by `allocproc`. The kernel portion of the process' page table, which converts virtual addresses above `KERNBASE` to physical addresses between 0 and `PHYSTOP`, is configured by `setupkvm`.

`proc.c`:

```
932 void create_kernel_process(const char *name, void (*entrypoint)()){
933
934     struct proc *p = allocproc();
935     if(p == 0) panic("create_kernel_process failed");
936     if((p->pgdir = setupkvm()) == 0) panic("setupkvm failed");
937
938     //This is a kernel process. Trap frame stores user space registers. We
don't need to initialise tf. Also, since this doesn't need to have a
userspace, we don't need to assign a size to this process. eip stores
address of next instruction to be executed
939     p->context->eip = (uint)entrypoint;
940     safestrcpy(p->name, name, sizeof(p->name));
941
942     acquire(&ptable.lock);
943     p->state = RUNNABLE;
944     release(&ptable.lock);
945
946 }
```

Task 2

There are several elements to this task. We must first create a process queue to keep track of the processes that were denied more memory because there were no vacant pages. A circular queue structure named `rq` was developed. And `rqueue` is the particular queue that contains processes with swap out requests. Additionally, we developed the `rq`-related methods `rpush()` and `rpop()`. We need to use a lock that has been initialised in `pinit` to access the queue. Additionally, we set the starting values of `s` and `e` in `userinit` to zero.

We also introduced prototypes in `defs.h` since the queue and the functions related to it are required in other files.

`proc.c`:

```
16 struct rq{
void
userinit(void)
{
    acquire(&rqueue.lock);
    rqueue.s=0;
    rqueue.e=0;
    release(&rqueue.lock);

    acquire(&rqueue2.lock);
    rqueue2.s=0;
    rqueue2.e=0;
    release(&rqueue2.lock);

    struct proc *p;
    (rqueue.e)%=NPROC;
    release(&rqueue.lock);
    return 1;
}
```

```
struct rq rqueue, rqueue2;
struct proc* rpop(){
    acquire(&rqueue.lock);
    if(rqueue.s==rqueue.e){
        release(&rqueue.lock);
    }
}

void
pinit(void)
{
    initlock(&ptable.lock, "ptable");
    initlock(&rqueue2.lock, "rqueue2");
    initlock(&rqueue.lock, "rqueue");
    initlock(&sleeping_channel_lock, "sleeping_channel");
}
```

`defs.h`:

```
void yield(void);
void create_kernel_process(const char *name, void (*entrypoint)());
void swap_out_process_function();
void swap_in_process_function();
extern int swap_out_process_exists;
extern int swap_in_process_exists;
extern struct rq rqueue;
extern struct rq rqueue2;
int rpush(struct proc *p);
int rpush2(struct proc *p);
struct proc* rpop();
struct proc* rpop2();
```

```
12 struct rq;
```

Now, `kalloc` returns 0 every time it is unable to allot pages to a process. This informs `allocuvn` that no memory was allocated for the required amount (`mem=0`). In this case, we must first set the process status to sleeping. The process

then has to be added to the queue for swap out requests (*Note: The process sleeps on a

unique sleeping channel named sleeping channel, which is protected by a lock called sleeping channel lock. sleeping channel count is used for exceptional instances when the system boots.)

vm.c:

```
14 char * sleeping_channel;
15 struct spinlock sleeping_channel_lock;
16 int sleeping_channel_count=0;
17
5 return 0;
6
7 a = PGROUNDUP(oldsz);
8 for(; a < newsz; a += PGSIZE){
9     mem = kalloc();
10    if(mem == 0){
11        // cprintf("allocvm out of memory\n");
12        deallocvm(pgdir, newsz, oldsz);
13        //SLEEP
14        myproc()->state=SLEEPING;
15        acquire(&sleeping_channel_lock);
16        myproc()->chan=sleeping_channel;
17        sleeping_channel_count++;
18        release(&sleeping_channel_lock);
19        rpush(myproc());
20        if(!swap_out_process_exists){
21            swap_out_process_exists=1;
22
23        create_kernel_process("swap_out_process",&swap_out_process_function);
24    }
25
26    return 0;
27 }
28 memset(mem, 0, PGSIZE);
```

allocvm:

Now, kalloc always returns 0 if a process cannot receive a page allocation. Allocvm is informed that the required memory wasn't allocated (mem=0) by this. We must first set the process status in this case to sleeping. (*Note: The process sleeps on a unique channel named sleeping channel, which is protected by a lock called sleeping channel lock; sleeping channel count is used for exceptional circumstances when the system boots.) Next, we must add the current process to the queue for swap out requests.

Next, we develop a system that wakes up all processes sleeping on the sleeping channel whenever free pages become available. In kalloc.c, we change kfree as follows:


```

60 void
61 kfree(char *v)
62 {
63     struct run *r;
64
65     if((uint)v % PGSIZE || v < end || V2P(v) >= PHYSTOP)
66         panic("kfree");
67
68     // Fill with junk to catch dangling refs.
69     //memset(v, 1, PGSIZE);
70     for(int i=0; i<PGSIZE; i++) v[i]= 1;
71
72     if(kmem.use_lock) acquire(&kmem.lock);
73     r = (struct run*)v;
74     r->next = kmem.freelist;
75     kmem.freelist = r;
76     if(kmem.use_lock) release(&kmem.lock);
77     // wake up process sleeping on sleeping channel
78     if(kmem.use_lock) acquire(&sleeping_channel_lock);
79     if(sleeping_channel_count){
80         wakeup(sleeping_channel);
81         sleeping_channel_count = 0;
82     }
83     if(kmem.use_lock) release(&sleeping_channel_lock);
84 }
85

```

In essence, all processes that were sent to sleep on the sleeping channel after being preempted because there weren't any pages available were. The wakeup() system function is used to wake up all sleeping processes on sleeping channel.

We'll now describe the changing out procedure. the function's swap out process entry point, which starts the switching out process. Due to the length of the function, We've included multiple screenshots:

screenshots:

```

235 void swap_out_process_function(){
236     acquire(&rqueue.lock);
237     while(rqueue.s!=rqueue.e){
238         struct proc *p=rpop();
239         pde_t* pd = p->pgdir;
240         for(int i=0; i<NPENTRIES; i++){
241             if(pd[i]&PTE_A) continue;
242             pte_t *pgtab = (pte_t*)P2V(PTE_ADDR(pd[i]));
243             for(int j=0; j<NPENTRIES; j++){
244                 //skip if found
245                 if((pgtab[j]&PTE_A) || !(pgtab[j]&PTE_P)) continue;
246                 pte_t *pte=(pte_t*)P2V(PTE_ADDR(pgtab[j]));
247
248                 int pid=p->pid;
249                 int virt = ((1<<22)*i)+((1<<12)*j);
250                 char c[50];
251
252                 int_to_string(pid,c);
253                 int x=strlen(c);
254                 c[x]='_';
255
256                 int_to_string(virt,c+x+1);
257
258
259

```

```

258
259     int_to_string(virt,c+x+1);
260     safestrcpy(c+strlen(c),".swp",5);
261
262     // file management
263     int fd=proc_open(c, O_CREATE | O_RDWR);
264     if(fd<0){
265
266         cprintf("error creating or opening file: %s\n", c);
267         panic("swap_out_process");
268     }
269
270
271
272
273     if(proc_write(fd,(char *)pte, PGSIZE) != PGSIZE){
274         cprintf("error writing to file: %s\n", c);
275         panic("swap_out_process");
276     }
277
278
279     proc_close(fd);
280
281     kfree((char*)pte);
282     memset(&pgtab[j],0,sizeof(pgtab[j]));

```

```

281     kfree((char*)pte);
282     memset(&pgtab[j],0,sizeof(pgtab[j]));
283
284     pgtab[j]=((pgtab[j])^(0x080));
285     break;
286 }
287 }
288
289 }
290
291 release(&rqueue.lock);
292
293 struct proc *p;
294 if((p=myproc())==0) panic("swap out process");
295
296 swap_out_process_exists=0;
297
298 p->killed = 0;
299 p->state = UNUSED;
300 p->parent = 0;
301 p->name[0] = '*';
302 sched();
303 }

```

Until the swap out requests queue (rqueue1) is not empty, the procedure iterates in a loop. A sequence of instructions are carried out to end the swap out process when the queue is empty (Image 2). The loop begins by removing the first process from the rqueue, then it searches its page table for a victim page using the LRU policy. We repeatedly go over each entry in the page table (pgdir) for the process and for every secondary page table,

retrieves the physical address. The accessed bit (A), which is the sixth bit from the right on each entry in the page table for each secondary page table, is examined throughout each iteration of the page table. By examining the bitwise & of the entry, we may determine if it is set and PTE_A (which we defined as 32 in mmu.c)).

Important information on the Accessible flag: When the scheduler switches the context of a process, all accessed bits are unset. As a result of our actions, the accessed bit detected by the swap out process function will show whether the item was accessed during the previous iteration of the process:

```
for(int i=0;i<NPENTRIES;i++){
    if(((p->pgdir)[i]&PTE_P && ((p->pgdir)[i]&PTE_A){
        pte_t* pgtab = (pte_t*)P2V(PTE_ADDR((p->pgdir)[i]));
        for(int j=0;j<NPENTRIES;j++){
            if(pgtab[j]&PTE_A){
                pgtab[j]^=PTE_A;
            }
        }
        ((p->pgdir)[i])^=PTE_A;
    }
}

// Switch to chosen process. It is the process's job
// to release ptable.lock and then reacquire it
// before jumping back to us.
c->proc = p;
switchvm(p);
p->state = RUNNING;

swtch(&(c->scheduler), p->context);
```

This function, which is included in the scheduler, essentially resets all bits that have been accessed in the process' primary and secondary page tables.

Returning to the swap out process function now. The victim page is selected by the function (using the macros stated in section A report) as soon as it locates a secondary page table entry with the accessed bit unset. The following step is to replace and store this page on drive.

To name the file that holds this page, we utilise the process' process id (pid, line 267 in the picture), as well as the virtual address of the page that has to be deleted (virt, line 268 in the image). A new function we call "int to string" converts an integer into a specified string. We create the filename using the numbers pid and virt using this method. This is the declaration for the function in proc.c:

```
26 void int_to_string(int x, char *c){
27     if(x==0){
28         c[0]='0';
29         c[1]='\0';
30         return;
31     }
32     int i=0;
33
34     while(x>0){
35         c[i]=x%10+'0';
36         i++;
37         x/=10;
38     }
39     c[i]='\0';
40
41     for(int j=0;j<i/2;j++){
42         char a=c[j];
43         c[j]=c[i-j-1];
44         c[i-j-1]=a;
45     }
46
47 }
```

The victim page's contents must be written to a file called `_.swp`. But this is where we get into trouble. The filename is kept in a string named `c`. Proc.c cannot make calls to file system functions.

The answer was to copy the open, write, read, close, etc. functions from `sysfile.c` to `proc.c`, modify them because `sysfile.c` methods took arguments in a different way, and then rename them to `proc open`, `proc read`, `proc write`, `proc close`, etc. so we could use them in `proc.c`.

Now, we write a page back to storage utilising these functions. We open a file with the permissions O CREATE and O RDWR (using proc open) (we have imported fcntl.h with these macros).

If this file doesn't already exist, O CREATE creates it, and O RDWR stands for read/write. The integer fd contains the file descriptor. We use proc write to write the page to this file using this file descriptor. Then, this page is put to the free page queue using kfree so that it is ready for usage (remember, when kfree adds a page to the free queue, we also wake up any processes sleeping on sleeping channel). The page table entry is then cleared using memset.

The next step is crucial since Task 3 requires us to determine whether or not the page that triggered a problem was replaced. We set the secondary page table entry's eighth rightmost bit (27) to indicate that this page has been switched out. To do this work, we use xor.

The loop is broken and the procedure is put on hold when the

queue is empty. We can't wipe the kstack of the ongoing kernel processes while they are being terminated since they won't know which process to execute next as a result. Their kstack has to be cleared from outside the process. To do this, we preempt the process first and then wait for the scheduler to locate it. The scheduler clears the kstack and name of any kernel processes it discovers that are in the UNUSED state. By looking at the name of the kernel process, which had its initial character altered to the symbol * when the process stopped, the scheduler may determine if it is in an idle state. Thus the ending of kernel processes has two parts:

```
49 int proc_write(int fd, char *p, int n){
50     struct file *f;
51     if(fd >= NOFILE || fd < 0 || (f=myproc()->ofile[fd]) == 0) return -1;
52     return filewrite(f, p, n);
53 }
54
55
56 int proc_close(int fd){
57     struct file *f;
58     if(fd < 0 || fd >= NOFILE || (f=myproc()->ofile[fd]) == 0) return -1;
59
60
61     myproc()->ofile[fd] = 0;
62     fileclose(f);
63     return 0;
64 }
65
66
67 int proc_read(int fd, int n, char *p){
68     struct file *f;
69     if(fd >= NOFILE || fd < 0 || (f=myproc()->ofile[fd]) == 0) return -1;
70     return fileread(f, p, n);
71
72 }
```

from within process

```
292
293 struct proc *p;
294 if((p=myproc())==0) panic("swap out process");
295
296 swap_out_process_exists=0;
297
298 p->killed = 0;
299 p->state = UNUSED;
300 p->parent = 0;
301 p->name[0] = '*';
302 sched();
303 }
```

from scheduler

```
706 if(p->state==UNUSED && p->name[0]=='*'){
707
708     kfree(p->kstack);
709     p->pid=0;
710     p->kstack=0;
711     p->name[0]=0;
712 }
713
714 if(p->state != RUNNABLE)
715     continue;
```

Task 3

First, a swap in request queue has to be created. To make a swap in the request queue known as rqueue2 in proc.c, we utilised the same struct (rq) as in Task 2. In defs.h, we additionally specify an extern prototype for rqueue2. We also developed the corresponding functions for rqueue2 (rpop2() and rpush2()) and stated their prototype in defs.h along with declaring the queue. Additionally, we initialised pinit's lock. Additionally, we set its s and e variables to zero in userinit.

Since all of the variables and functions are similar to those in Task 2, I won't attach screenshots of them here.

Next, we add a new entry named addr to the struct proc in proc.h (int). The swapping in function will be informed by this item of the virtual address where the page fault occurred.

proc.h

```
1 char name[16];
2 int addr;
3 };
4
```

The page fault (T_PGFLT) traps generated by trap.c must then be handled. In a function named handlePageFault(), we carry it out.

trap.c

```
20 void handlePageFault(){
21     struct proc *p=myproc();
22     int addr=rcr2();
23     acquire(&swap_in_lock);
24     sleep(p,&swap_in_lock);
25     pde_t *pde = &(p->pgdir)[PDX(addr)];
26     pte_t *pgtab = (pte_t*)P2V(PTE_ADDR(*pde));
27
28     if((pgtab[PDX(addr)]&0x000){
29         p->addr = addr;
30         rpush2(p);
31         if(!swap_in_process_exists){
32             swap_in_process_exists=1;
33             create_kernel_process("swap_in_process", &swap_in_process_function);
34         }
35     }
36     else {
37         // ...
38     }
39 }
```

```
103 lapiceoi();
104 break;
105 case T_PGFLT:
106     handlePageFault();
107 break;
108 //PAGEBREAK: 13
109 default:
```

Similar to Part A, `handlePageFault` uses `rcr2` to identify the virtual address at which the page fault occurred (). We then use a brand-new lock called the swap in lock to put the active process to sleep (initialised in `trap.c` and with `extern` in `defs.h`). The page table entry corresponding to this address is then obtained (the logic is identical to `walkpgdir`). We must now determine if this page was switched out. When switching out a page in Task 2, we set the page table entry's bit of the seventh order (27). Beginning on the fifth page of this report, this is discussed.

So, using bitwise `&` and `0x080`, we examine the page's 7th order bit to determine whether or not the page was switched out. Initiating swap in process (assuming it doesn't already exist - verify using `swap in process` exists) is done if it is set. Otherwise, as instructed by the assignment, we may safely halt the operation using `exit()`.

We now proceed with the switching in procedure. As you can see in `handlePageFault`, the entry point for the swapping out process is `swap in process` function (defined in `proc.c`).

Due to its length, `swap in process` function is displayed on the next page. For the real function, see the page after this one.

In the Task 2 section of the report, We have previously discussed how We introduced file management features in `proc.c`. Here, We will only briefly discuss the functions We used and how We used them. The loop in the function continues until `rqueue2` is not empty. In the loop, a process is selected from the queue, and its `pid` and `addr` values are extracted to get the file name. The filename is then created as a string named `"c"` using the `int` to string function. Then, it used `proc open` to open this file using file descriptor `fd` in read-only mode (`O_RDONLY`). Then, using `kalloc`, we provide this process access to a free frame (`mem`). Using `proc read`, we read into this free frame data from the file with the `fd` file descriptor. Then, after removing the static keyword from it in `vm.c`, we make `mappages` accessible to `proc.c` by declaring a prototype there. The physical page that was obtained using `kalloc` and read into is then mapped using `mappages` to the page corresponding to `addr` (`mem`).

Then, using `wakeup` to correct the page fault, we wake up the process for which a new page was allocated. We execute the kernel process termination instructions when the loop has finished.

```

24 int mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm);

307 void swap_in_process_function(){
308     acquire(&rqueue2.lock);
309     while(rqueue2.s!=rqueue2.e){
310         struct proc *p=rpop2();
311         int pid=p->pid;
312         int virt=PTE_ADDR(p->addr);
313
314         char c[50];
315         int_to_string(pid,c);
316         int x=strlen(c);
317         c[x]='_';
318
319         int_to_string(virt,c+x+1);
320         safestrcpy(c+strlen(c),".swp",5);
321         int fd=proc_open(c,O_RDONLY);
322         if(fd<0){
323             release(&rqueue2.lock);
324             fprintf("could not find page file in memory: %s\n",
325                 c);
326             panic("swap_in_process");
327         }
328     }
329 }

331
332     proc_read(fd,PGSIZE,mem);
333
334     if(mappages(p->pgdir, (void *)virt, PGSIZE, V2P(mem), PTE_W|
335         PTE_U)<0){
336         release(&rqueue2.lock);
337         panic("mappages");
338     }
339     wakeup(p);
340 }
341 release(&rqueue2.lock);
342 struct proc *p;
343 if((p=myproc())==0)panic("swap_in_process");
344
345 swap_in_process_exists=0;
346 p->killed = 0;
347
348
349 p->state = UNUSED;
350 p->parent = 0;
351
352
353 p->name[0] = '*';
354 sched();

```

Task 4

In this section, we'll build a testing framework to evaluate the functions that we developed in the other sections. We'll utilise a user-space application called memtest to carry out this task on our behalf. Below is a description of how memtest is implemented.

By examining the implementation, we may draw the following conclusions:

- Using the fork() system function, the primary process generates 20 child processes.
- Every child process runs a 10 iterations long loop.

```
1 #include "types.h"
2 #include "stat.h"
3 #include "user.h"
4
5
6 int main(int argc, char* argv){
7     for(int i=0; i<20; i++){
8         if(!fork()){
9             printf(1, "Child %d\n", i+1);
10            printf(1, "Iteration Matched Different\n");
11            int j = 0;
12            while(j<10){
13                int *arr = malloc(4096);
14                for(int k=0; k<1024; k++) arr[k] = k*k - 4*k + 1;
15                int matched=0;
16                for(int k=0; k<1024; k++){
17                    if(arr[k] == k*k - 4*k + 1) matched+=4;
18                }
19                if(j<9) printf(1, "    %d    %dB\n",
20                    j+1, matched, 4096-matched);
21                else printf(1, "    %d    %dB    %dB\n",
22                    j+1, matched, 4096-matched);
23                j++;
24            }
25        }
26    }
27    while(wait()!= -1);
28    exit();
29 }
```

```
23
24
25
26
27     printf(1, "\n");
28
29     exit();
30 }
31
32
33
34
35
36 while(wait()!= -1);
37 exit();
38
39 }
```

The result shows that our implementation passes the sanity test since every index stores the right value.

We now run the tests on various PHYSTOP values to better evaluate our implementation (defined in memlayout.h). 0xE0000000 is the default value for

- Using malloc, 4096B (4KB) of memory is allocated per loop ()

- The mathematical equation that is calculated using math func provides the value stored at array index We ().

- The number of bytes that have the correct values is kept in a counter called matched. This is accomplished by comparing the value provided by the function for each index with the value saved at each position.

We must add memtest to the Makefile's UPROGS and EXTRA sections so that the xv6 user can access it in order to execute it.

We get the following output from memtest.

```
$ memtest
Child 1
Iteration Matched Different
1      4096B      0B
2      4096B      0B
3      4096B      0B
4      4096B      0B
5      4096B      0B
6      4096B      0B
7      4096B      0B
8      4096B      0B
9      4096B      0B
10     4096B      0B
```

```
Child 2
Iteration Matched Different
1      4096B      0B
2      4096B      0B
3      4096B      0B
```

PHYSTOP (224MB). Its value was modified to 0x0400000 (4MB). We choose 4MB since xv6 requires this amount of RAM in order to run kinit1. The output from memtest was similar to the prior output, proving that the implementation is sound.